

6.824 2015 Lecture 1: Introduction and lab overview

6.824: Distributed Systems Engineering

What is a distributed system?

- multiple networked cooperating computers

Examples: Internet E-Mail, Athena file server, Google MapReduce, etc.

Why distribute?

- to connect physically separate entities
- to achieve security via physical isolation
- to tolerate faults via replication at separate sites
- to increase performance via parallel CPUs/mem/disk/net

But:

- complex, hard to debug
- new classes of problems, e.g. partial failure (did server accept my e-mail?)
- advice: don't distribute if a central system will work

Why take this course?

- interesting -- hard problems, non-obvious solutions
- active research area -- lots of progress + big unsolved problems
- used by real systems -- driven by the rise of big Web sites
- hands-on -- you'll build a real system in the labs

COURSE STRUCTURE

<http://pdos.csail.mit.edu/6.824>

Course components:

Lectures about big ideas, papers, labs

Readings: research papers as case studies

- please read papers before class
 - otherwise boring, and you can't pick it up by listening
- each paper has a question for you to answer
- and you should think of a question you would like to have answered
- submit question&answer before class, one or two paragraphs

Mid-term exam in class, and final exam

Labs: build increasingly sophisticated fault-tolerant services

First lab is due on Monday

For PhD students, you can substitute a small research project for 5th lab

talk to us

TAs: Steven Allen, Rohan Mahajan, Steven Valdez

- answer questions about material
- help you with labs
- will post office hours

MAIN TOPICS

Example:

- a shared file system, so users can cooperate, like Athena's AFS
- lots of client computers
- [diagram: clients, network, vague set of servers]

Topic: architecture

What interface?

- Clients talk to servers -- what do they say?
- File system (files, file names, directories, &c)?

- Disk blocks, with FS in client?
- Separate naming + file servers?
- Separate FS + block servers?
- Single machine room or unified wide area system?
- Wide-area more difficult.
- Transparent?
 - i.e. should it act exactly like a local disk file system?
 - or is it OK if apps/users have to cope with distribution,
 - e.g. know what server files are on, or deal with failures.
- Client/server or peer-to-peer?
- All these interact w/ performance, usefulness, fault behavior.

Topic: implementation

- How to simplify network communication?
 - Can be messy (msg formatting, re-transmission, host names, &c)
 - Frameworks can help: RPC, MapReduce, &c
- How to cope with inherent concurrency?
 - Threads, locks, &c.

Topic: performance

- Distribution can hurt: network b/w and latency bottlenecks
 - Lots of tricks, e.g. caching, concurrency, pre-fetch
- Distribution can help: parallelism, pick server near client
- Idea: scalable design
 - Nx servers -> Nx total performance
- Need a way to divide the load by N
 - Divide data over many servers ("sharding" or "partitioning")
 - By hash of file name?
 - By user?
 - Move files around dynamically to even out load?
 - "Stripe" each file's blocks over the servers?
- Performance scaling is rarely perfect
 - Some operations are global and hit all servers (e.g. search)
 - Nx servers -> 1x performance
 - Load imbalance
 - Everyone wants to get at a single popular file
 - > one server 100%, added servers mostly idle
 - > Nx servers -> 1x performance

Topic: fault tolerance

- Big system (1000s of server, complex net) -> always something broken
- We might want:
 - Availability -- I can keep using my files despite failures
 - Durability -- my files will come back to life someday
- Availability idea: replicate
 - Servers form pairs, each file on both servers in the pair
 - Client sends every operation to both
 - If one server down, client can proceed using the other
- Opportunity: operate from both "replicas" independently if partitioned?
- Opportunity: can 2 servers yield 2x availability AND 2x performance?

Topic: consistency

- Assume a contract w/ apps/users about meaning of operations
 - e.g. "read yields most recently written value"
- Consistency is about fulfilling the contract
 - despite failure, replication/caching, concurrency, &c
- Problem: keep replicas identical
 - If one is down, it will miss operations
 - Must be brought up to date after reboot
 - If net is broken, *both* replicas maybe live, and see different ops
 - Delete file, still visible via other replica
 - "split brain" -- usually bad
- Problem: clients may see updates in different orders
 - Due to caching or replication
 - I make 6.824 directory private, then TA creates grades file

LABS

3/6

```

input is thousands of text files
Map(k, v)
    split v into words
    for each word w
        emit(w, "1")
Reduce(k, v)
    emit(len(v))

```

What does MR framework do for word count?

```

[master, input files, map workers, map output, reduce workers, output files]
input files:
    f1: a b
    f2: b c
send "f1" to map worker 1
    Map("f1", "a b") -> <a 1> <b 1>
send "f2" to map worker 2
    Map("f2", "b c") -> <b 1> <c 1>
framework waits for Map jobs to finish
workers sort Map output by key
framework tells each reduce worker what key to reduce
    worker 1: a
    worker 2: b
    worker 2: c
each reduce worker pulls needed Map output from Map workers
    worker 1 pulls "a" Map output from every worker
each reduce worker calls Reduce once for each of its keys
    worker 1: Reduce("a", [1]) -> 1
    worker 2: Reduce("b", [1, 1]) -> 2
               Reduce("c", [1]) -> 1

```

Why is the MR framework convenient?

- * programmer only needs to think about the core work, the Map and Reduce functions, does not have to worry network communication, failure, &c.
- * the grouping by key between Map and Reduce fits some applications well (e.g., word count), since it brings together data needed by the Reduce.
- * but some applications don't fit well, because MR only allows the one type of communication between different parts of the application.
e.g. word count but sort by frequency.

Why might MR have good performance?

```

Map and Reduce functions run in parallel on different workers
Nx workers -> divide run-time by N
But rarely quite that good:
    move map output to reduce workers
    stragglers
    read/write network file system

```

What about failures?

```

People use MR with 1000s of workers and vast inputs
Suppose each worker only crashes once per year
    That's 3 per day!
So a big MR job is very likely to suffer worker failures
Other things can go wrong:
    Worker may be slow
    Worker CPU may compute incorrectly
    Master may crash
    Parts of the network may fail, lose packets, &c
    Map or Reduce or framework may have bugs in software

```

Tools for dealing with failure?

```

retry -- if worker fails, run its work on another worker
replicate -- run each Map and Reduce on *two* workers

```

replace -- for long-term health
 MapReduce uses all of these

Puzzles for retry

how do we know when to retry?
 can we detect when Map or Reduce worker is broken?
 can we detect incorrect worker output?
 can we distinguish worker failure from worker up, network lossy?
 why is retry correct?
 what if Map produces some output, then crashes?
 will we get duplicate output?
 what if we end up with two of the same Map running?
 in general, calling a function twice is not the same as calling it once
 why is it OK for Map and Reduce?

Helpful assumptions

One must make assumptions, otherwise too hard
 No bugs in software
 No incorrect computation: worker either produces correct output,
 or nothing -- assuming fail-stop.
 Master doesn't crash
 Map and Reduce are strict functions of their arguments
 they don't secretly read/write files, talk to each other,
 send/receive network messages, &c

lab 1 has three parts:

Part I: just Map() and Reduce() for word count
 Part II: we give you most of a distributed multi-server framework,
 you fill in the master code that hands out the work
 to a set of worker threads.
 Part III: make master cope with crashed workers by re-trying.

Part I: main/wc.go

stubs for Map and Reduce
 you fill them out to implement word count
 Map argument is a string, a big chunk of the input file

demo of solution to Part I

```
./wc master kjvl2.txt sequential
more mrtmp.kjvl2.txt-1-2
more mrtmp.kjvl2.txt
```

Part I sequential framework: mapreduce/mapreduce.go RunSingle()
 split, maps, reduces, merge

Part II parallel framework:

master
 workers...
 shared file system
 our code splits the input before calling your master,
 and merges the output after your master returns
 our code only tells the master the number of map and reduce splits (jobs)
 each worker sends Register RPC to master
 your master code must maintain a list of registered workers
 master sends DoJob RPCs to workers
 if 10 map jobs and 3 workers,
 send out 3, wait until one worker says it's done,
 send it another, until all 10 done
 then the same for reduces
 master only needs to send job # and map vs reduce to worker
 worker reads input from files
 so your master code only needs to know the number of
 map and reduce jobs!
 which it can find from the "mr" argument

Thursday:

- master and workers talk via RPC, which hides network complexity
- more about RPC on Thursday