

## 6.824 2015 Lecture 2: Infrastructure: RPC and threads

## Remote Procedure Call (RPC)

a key piece of distrib sys machinery; all the labs use RPC  
 goal: easy-to-program network communication  
 hides most details of client/server communication  
 client call is much like ordinary procedure call  
 server handlers are much like ordinary procedures  
 RPC is widely used!

RPC ideally makes net communication look just like fn call:

Client:

z = fn(x, y)

Server:

```
fn(x, y) {
  compute
  return z
}
```

RPC aims for this level of transparency

Examples from lab 1:

DoJob

Register

RPC message diagram:

```
Client          Server
request--->
<---response
```

Software structure

```
client app      handlers
stubs           dispatcher
RPC lib         RPC lib
net  -----  net
```

A few details:

Which server function (handler) to call?

Marshalling: format data into packets

Tricky for arrays, pointers, objects, &c

Go's RPC library is pretty powerful!

some things you cannot pass: e.g., channels, functions

Binding: how does client know who to talk to?

Maybe client supplies server host name

Maybe a name service maps service names to best server host

Threads:

Client often has many threads, so >1 call outstanding, match up replies

Handlers may be slow, so server often runs each in a thread

RPC problem: what to do about failures?

e.g. lost packet, broken network, slow server, crashed server

What does a failure look like to the client RPC library?

Client never sees a response from the server

Client does *\*not\** know if the server saw the request!

Maybe server/net failed just before sending reply  
 [diagram of lost reply]

Simplest scheme: "at least once" behavior

RPC library waits for response for a while

If none arrives, re-send the request

Do this a few times

Still no response -- return an error to the application

Q: is "at least once" easy for applications to cope with?

Simple problem w/ at least once:

client sends "deduct \$10 from bank account"

Q: what can go wrong with this client program?

Put("k", 10) — an RPC to set key's value in a DB server

Put("k", 20) — client then does a 2nd Put to same key

[diagram, timeout, re-send, original arrives very late]

Q: is at-least-once ever OK?

yes: if it's OK to repeat operations, e.g. read-only op

yes: if application has its own plan for coping w/ duplicates  
which you will need for Lab 1

Better RPC behavior: "at most once"

idea: server RPC code detects duplicate requests

returns previous reply instead of re-running handler

Q: how to detect a duplicate request?

client includes unique ID (XID) with each request

uses same XID for re-send

server:

```
if seen[xid]:
```

```
    r = old[xid]
```

```
else
```

```
    r = handler()
```

```
    old[xid] = r
```

```
    seen[xid] = true
```

some at-most-once complexities

this will come up in labs 2 and on

how to ensure XID is unique?

big random number?

combine unique client ID (ip address?) with sequence #?

server must eventually discard info about old RPCs

when is discard safe?

idea:

unique client IDs

per-client RPC sequence numbers

client includes "seen all replies <= X" with every RPC

much like TCP sequence #s and acks

or only allow client one outstanding RPC at a time

arrival of seq+1 allows server to discard all <= seq

or client agrees to keep retrying for < 5 minutes

server discards after 5+ minutes

how to handle dup req while original is still executing?

server doesn't know reply yet; don't want to run twice

idea: "pending" flag per executing RPC; wait or ignore

What if an at-most-once server crashes and re-starts?

if at-most-once duplicate info in memory, server will forget

and accept duplicate requests after re-start

maybe it should write the duplicate info to disk?

maybe replica server should also replicate duplicate info?

What about "exactly once"?

at-most-once plus unbounded retries plus fault-tolerant service

Lab 3

Go RPC is "at-most-once"

open TCP connection

write request to TCP connection

TCP may retransmit, but server's TCP will filter out duplicates

no retry in Go code (i.e. will NOT create 2nd TCP connection)

Go RPC code returns an error if it doesn't get a reply

perhaps after a timeout (from TCP)

perhaps server didn't see request  
 perhaps server processed request but server/net failed before reply came back

Go RPC's at-most-once isn't enough for Lab 1  
 it only applies to a single RPC call  
 if worker doesn't respond, the master re-send to it to another worker  
 but original worker may have not failed, and is working on it too  
 Go RPC can't detect this kind of duplicate  
 No problem in lab 1, which handles at application level  
 Lab 2 will explicitly detect duplicates

## Threads

threads are a fundamental server structuring tool  
 you'll use them a lot in the labs  
 they can be tricky  
 useful with RPC  
 Go calls them goroutines; everyone else calls them threads

Thread = "thread of control"

threads allow one program to (logically) do many things at once  
 the threads share memory  
 each thread includes some per-thread state:  
 program counter, registers, stack

## Threading challenges:

sharing data  
 two threads modify the same variable at same time?  
 one thread reads data that another thread is changing?  
 these problems are often called races  
 need to protect invariants on shared data  
 use Go sync.Mutex

coordination between threads  
 e.g. wait for all Map threads to finish  
 use Go channels

deadlock  
 thread 1 is waiting for thread 2  
 thread 2 is waiting for thread 1  
 easy detectable (unlike races)

lock granularity  
 coarse-grained -> simple, but little concurrency/parallelism  
 fine-grained -> more concurrency, more races and deadlocks  
 let's look at a toy RPC package to illustrate these problems

look at today's handout -- l-rpc.go

it's a simplified RPC system  
 illustrates threads, mutexes, channels  
 it's a toy, though it does run  
 assumes connection already open  
 only supports an integer arg, integer reply  
 omits error checks

## struct ToyClient

client RPC state  
 mutex per ToyClient  
 connection to server (e.g. TCP socket)  
 xid -- unique ID per call, to match reply to caller  
 pending[] -- chan per thread waiting in Call()  
 so client knows what to do with each arriving reply

## Call

application calls reply := client.Call(procNum, arg)  
 procNum indicates what function to run on server  
 WriteRequest knows the format of an RPC msg  
 basically just the arguments turned into bits in a packet  
 Q: why the mutex in Call()? what does mu.Lock() do?

Q: could we move "xid := tc.xid" outside the critical section?  
 after all, we are not changing anything  
 [diagram to illustrate]

Q: do we need to WriteRequest inside the critical section?

note: Go says you are responsible for preventing concurrent map ops  
 that's one reason the update to pending is locked

Listener

runs as a background thread  
 what is <- doing?  
 not quite right that it may need to wait on chan for caller

Back to Call()...

Q: what if reply comes back very quickly?  
 could Listener() see reply before pending[xid] entry exists?  
 or before caller is waiting for channel?

Q: should we put reply:=<-done inside the critical section?  
 why is it OK outside? after all, two threads use it.

Q: why mutex per ToyClient, rather than single mutex per whole RPC pkg?

Server's Dispatcher()

note that the Dispatcher echos the xid back to the client  
 so that Listener knows which Call to wake up

Q: why run the handler in a separate thread?

Q: is it a problem that the dispatcher can reply out of order?

main()

note registering handler in handlers[]  
 what will the program print?

Q: when to use channels vs shared memory + locks?

here is my opinion

use channels when you want one thread to explicitly wait for another  
 often wait for a result, or wait for the next request  
 e.g. when client Call() waits for Listener()

use shared memory and locks when the threads are not intentionally  
 directly interacting, but just happen to r/w the same data  
 e.g. when Call() uses tc.xid

but: they are fundamentally equivalent; either can always be used.

Go's "memory model" requires explicit synchronization to communicate!

This code is not correct:

```
var x int
done := false
go func() { x = f(...); done = true }
while done == false { }
```

it's very tempting to write, but the Go spec says it's undefined  
 use a channel or sync.WaitGroup instead

Study the Go tutorials on goroutines and channels